

BEN-GURION UNIVERSITY OF THE NEGEV

ADVANCED SYSTEM PROGRAMMING

202.1.6331

---

## Content Addressable Filesystem

---

*Authors:*

Gregory Guterman (345956502) Ameen Kassem ()

Publish date: November 5th, 2025

Submission date: November 19th, 2025

# 1 Project Structure

```
asp-caf-assignment/
+- deployment/
|   +- Dockerfile           # Development environment setup
+- Makefile                 # Build and development commands
+- assignment/
+- caf/
|   +- pyproject.toml       # Python package configuration
|   +- caf/
|       +- __main__.py      # Entry point
|       +- cli.py            # Command-line interface
|       +- cli_commands.py   # Command implementations
+- libcaf/
|   +- CMakeLists.txt       # CMake build configuration
|   +- pyproject.toml       # Python package configuration
|   +- libcaf/
|       +- constants.py    # Constants and configuration
|       +- plumbing.py     # Low-level repo operations
|       +- ref.py           # Reference handling
|       +- repository.py   # Repository management and high-level API
+- src/
|   +- bind.cpp             # Python bindings
|   +- blob.h               # Blob object definitions
|   +- caf.cpp/h            # Low-level C++ implementation
|   +- commit.h              # Commit object definitions
|   +- hash_types.cpp/h     # Hashing implementations
|   +- object_io.cpp/h      # Object I/O operations
|   +- tree.h                # Tree object definitions
|   +- tree_record.h        # Tree record structures
+- tests/
    +- caf/                  # CLI tests
    +- libcaf/                # Core library tests
```

# 2 Project Description

1. DockerFile: This file is crucial for setup of the project, using this file Docker image is created, and hence we can use a unified operating system to run the project; moreover, we can use a Docker container to build an actual release product. Also, all the requirements are installed in Docker, so they keep your computer clean, and you have all the up-to-date versions of Python libraries.
2. MakeFile: This file is a “manager” of the project, to build a Docker, to run it, and to run pytests, we use a Makefile. It is not common to have a makefile for a Python project, but since we have part of the project in C++, it’s crucial to have one.
3. caf/caf/\_\_main\_\_.py: This file is an entry point to the project; it gives us an option to run caf as an installed command as well as a Python module. It has a Python mechanism to have a “cpp-like” main function using \_\_main\_\_. After entry, it calls the CLI module, which contains the main logic of the project.

4. `caf/caf/cli.py`: This file implements a command line interface on the hight level, actual command implementation are in `cli_commands.py` for the project, in it user commands are defined (`init`, `commit`, etc). The program utilize a `argparse` library to process command line input to the python. It uses a data structure to map command names to their functions and definitions.
5. `caf/caf/cli_commands.py`: This file contains the actual implementation of the command-line logic for the CAF, available to users. Each function in this module represents one command that the user can address through the CLI, such as initializing the repository, committing, creating/deleting branches ans etc. The program interacts with a `libcaf` library throught modules like `repository.py` and others. This file connects user commands to the internal functionality of the CAF system.
6. `libcaf/libcaf/plumbing.py`: This file contains low-level functions that interact directly with the core C++ code. It provides functionailty for reading and writing objects, files, hashing and etc, higher level modules rely on this module to perform their operations efficiently. In more simple words this file is like a bridge between C++ and Python parts of the project.
7. `libcaf/libcaf/ref.py`: This file defines different types of references in CAF. A reference maybe a direct hash to a commit, or symbolic reference that points to another reference such as branch name. The model aslo has read and write functions from files, also validating references formart. Essentialy, this part of the project is need to manage commits and branches within repository.
8. `libcaf/libcaf/repository.py`: This file is the core of the Python interface to the CAF repo. The Repository class handels all the methods repository has, such as initializtion, creating commits, showing difference and etc. It is a backend for CLI commands and uses plumbing, ref and C++ modules to perform its operations. This file is a main controller of the repository state and caf functionality.
9. `libcaf/src/bind.cpp`: This file contains Python bidings for the C++ code. It creates the `_libcaf` Python module and bind all the low-level C++ and data structures be accesible from Python, In practice, this is a bridge between Python and C++ backend.
10. `libcaf/src/caf.cpp`: This file implements the low-level content-addressable storage operation in C++. Here realized core functionality of hashing files, saving the objects store, reading, writing object data by hash also deleting stored information. It is connects to Python via `pybind11` in `bind.cpp`.
11. `libcaf/src/hash_types.cpp/h`: This file implements hashing logic for Blobs, Trees and Commits. It communicates with `caf.cpp` and assist it to make content addressable storage work.
12. `libcaf/src/object_io.cpp/h`: This file manages input and output operations, how commits and trees are stored on and loaded from the disk. It writes object in a binary format to a file, identified by repectful hash.
13. `libcaf/src/CMakeLists.txt`: This file is a build for C++ part of the project.
14. `test/caf/cli_commands`: This is a seriece of files that test the CLI commands, that user interracts with, using `pytest` framework.
15. `test/libcaf`: This is a seriece of files that test the `libcaf` library, in partuclar methods that are available to the user and crutial to the CAF workflow, using `pytest` framework.

### 3 Bugs fixes

Fixed tests 1 – 3 that were failing due to unsorted dictionary iteration in `repository.py` file.

Names of the test are:

- tests/libcaf/test\_diff.py::test\_diff\_nested\_trees — FAILED

This test creates two directories `dir1` and `dir2`. In the second commit, `file_a.txt` in `dir1` is modified, and `file_b.txt` in `dir2` is replaced by `file_c.txt`, so both directories are considered modified. The test expects the first entry in the `modified` list to be `dir1`, but the actual value was `dir2`. This happens because the diff tree children were appended in traversal order, since we don't have a guarantee to sort in lexicographical order by directory name.

To fix this, I added a step in `diff_commits` that recursively sorts `diff.children` at every node by `record.name`, such that `dir1` stored before `dir2` in the `modified` list.

[Link to Commit](#)

- tests/libcaf/test\_diff.py::test\_diff\_moved\_file\_added\_first — FAILED

This test also sets up `dir1` and `dir2`, but then moves `file_a.txt` from `dir1` into `dir2` as `file_c.txt`. As a result, both `dir1` and `dir2` are reported as modified. The test expects `modified[0].record.name` to be `dir1`, but it was `dir2` instead. Again, the problem was that the diff children were stored in the order they were discovered during tree traversal, without any explicit sorting, so the order of modified directories was not deterministic.

To fix this, I added a step in `diff_commits` that recursively sorts `diff.children` at every node by `record.name`, such that `dir1` stored before `dir2` in the `modified` list.

[Link to Commit](#)

- tests/libcaf/test\_diff.py::test\_diff\_moved\_file\_removed\_first — FAILED

Here the setup is similar, but `file_b.txt` is moved from `dir2` into `dir1` as `file_c.txt`. This again makes both `dir1` and `dir2` modified. The test asserts that the first modified directory is `dir1`, but `modified[0].record.name` was `dir2`. The issue is the same: the order of diff nodes reflected traversal order rather than a sorted order by directory name.

To fix this, I added a step in `diff_commits` that recursively sorts `diff.children` at every node by `record.name`, such that `dir1` stored before `dir2` in the `modified` list.

[Link to Commit](#)

```
tests/libcaf/test_diff.py::test_diff_nested_trees PASSED
tests/libcaf/test_diff.py::test_diff_moved_file_added_first PASSED
tests/libcaf/test_diff.py::test_diff_moved_file_removed_first PASSED
tests/libcaf/test_hashing.py::test_hash_file_non_existent_file PASSED
```

Figure 1: Prove that test passes now

Fixed test 4 that was failing due to wrong tree record storage in tree.h file.

tests/libcaf/test\_objects.py::test\_tree\_entries\_are\_canonicalized — FAILED

This test checks that tree values are stored in lexicographical order by their names, not by order they are added or any other random order. The test constructs three `Tree` objects with the same files `a_file`, `a_file`, `c_file` but in different orders. The test expects all three trees to be stored in alphabetical order. However, the actual implementation of `tree.h` used `unordered_map`, which does not preserve any order, we used `map` to preserve the order and hence we get a tree entries stored in order by name as the test expects.

[Link to Commit](#)

```
tests/libcaf/test_objects.py::test_tree_entries_are_canonicalized PASSED
```

Figure 2: Prove that test passes now

## 4 Tags

Git tag used to mark specific points in the repository history, it is a static, immutable reference to the commit. It's difference from a branch is that the branch is a mutable pointer that automatically moves forward with next commit. Tag permanently pointer to a specific commit hash and dont change. Tags also differ from just commits, tag is a name for this commit, not the commit itself. Most common use case of tag is to mark release versions of a product. Git stores tags in refs/tags directory, while lightweight tags are simple pointer of commit, annotated tags also store metadata, that makes them a choice for formal release.